

PostGIS



Geospatial Objects for PostgreSQL

Tips for the PostGIS Power User



Topics

- PostGIS functions
 - Geometry constructors / deconstructors
accessors / spatial predicates
 - Walk through a few examples.
- DE-9IM
 - Fine-tuning spatial predicates
- PostgreSQL
 - Table inheritance / partitioning
 - Database tuning

Introduction

- What is PostGIS?
 - A PostgreSQL database extension that "spatially enables" the server back-end to support the storage of geometric objects in an object-relational PostgreSQL database.
 - <http://postgis.refrations.net/docs/>

Introduction


- Geometry Constructors
 - ST_GeomFromText(text)
 - ST_GeomFromWKB(bytea)

```
SELECT ST_GeomFromText( 'POINT(1718098 616348)', 3005 );
```

- Creates a geometry data type in EPSG:3005 projection

Introduction

- Geometry Constructors



```
SELECT ST_MakeLine(gps_points.geom)
FROM (SELECT geom
      FROM gps_point_data
      WHERE time_stamp::date = '2007-09-22'::date
      ORDER BY time_stamp
      ) AS gps_points;
```


Introduction

- Geometry Constructors
 - ST_BuildArea()



```
SELECT ST_BuildArea(ST_Collect(geom))  
FROM ...
```

Introduction

- Geometry Accessors / Deconstructors
 - ST_StartPoint()
 - ST_PointN(geometry, int)
 - ST_ExteriorRing(geometry)

```
SELECT ST_StartPoint(geom)
FROM my_lines;
```

Introduction

- Geometry Accessors / Deconstructors

```
SELECT point
FROM (SELECT ST_StartPoint(geom) AS point
      FROM my_lines
      UNION ALL
      SELECT ST_EndPoint(geom) AS point
      FROM my_lines) AS a
GROUP BY point
HAVING count(*) = 4;
```



Caution: GROUP BY uses a geometry's bounding box

Introduction

- Geometry Accessors / Deconstructors

```
SELECT ST_GeometryN( geom, 1 ) FROM my_multilines;
```

- How to explode a MULTI* table

```
SELECT generate_series( 1, 5 );
```

```
generate_series
```

```
-----
```

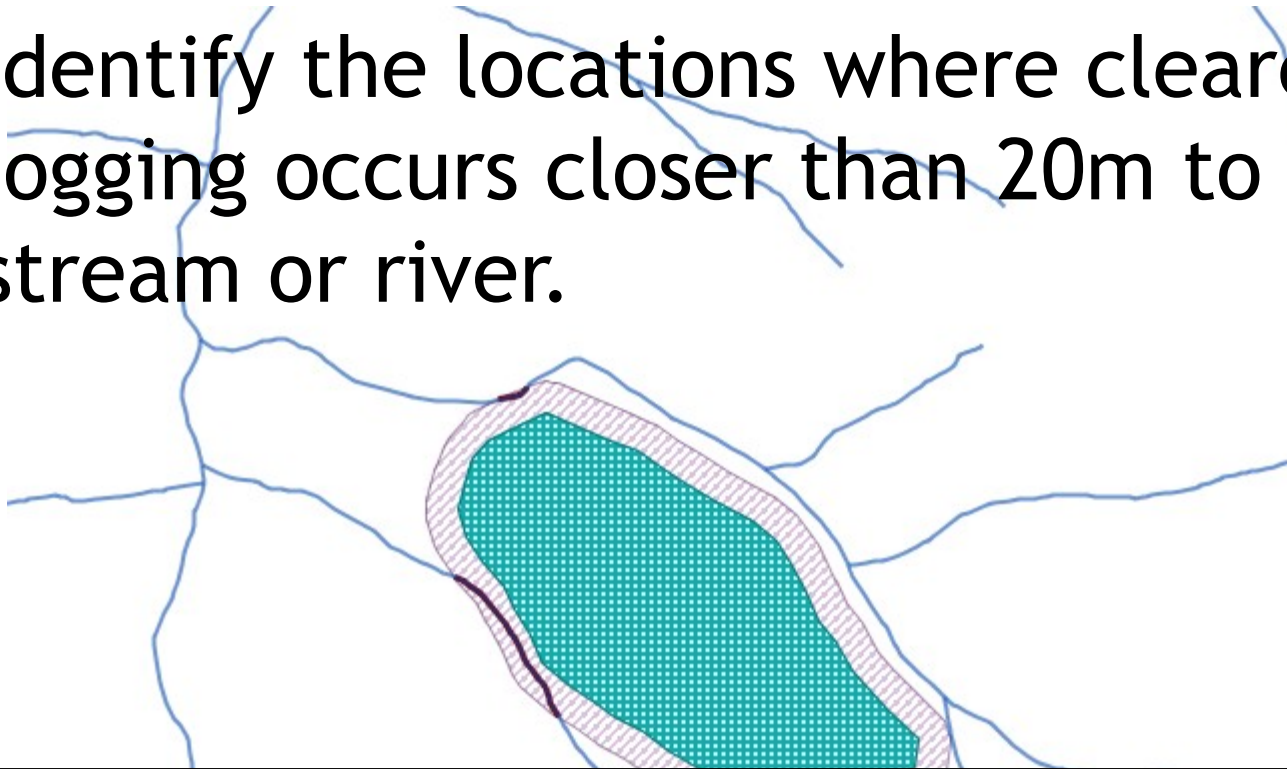
```
CREATE TABLE my_lines AS
SELECT geom
FROM (SELECT ST_GeometryN(geom,
                        generate_series(1, ST_NumGeometries(geom)))
      AS geom
FROM my_multilines
) AS foo;
```

Introduction

- Geometry Spatial Predicates / Functions
 - ST_Intersects()
 - ST_Within()
 - ST_Touches()
 - ST_GeomUnion()
 - ST_SymmetricDifference()
 - ST_ConvexHull()
 - ...

Sample PostGIS Queries

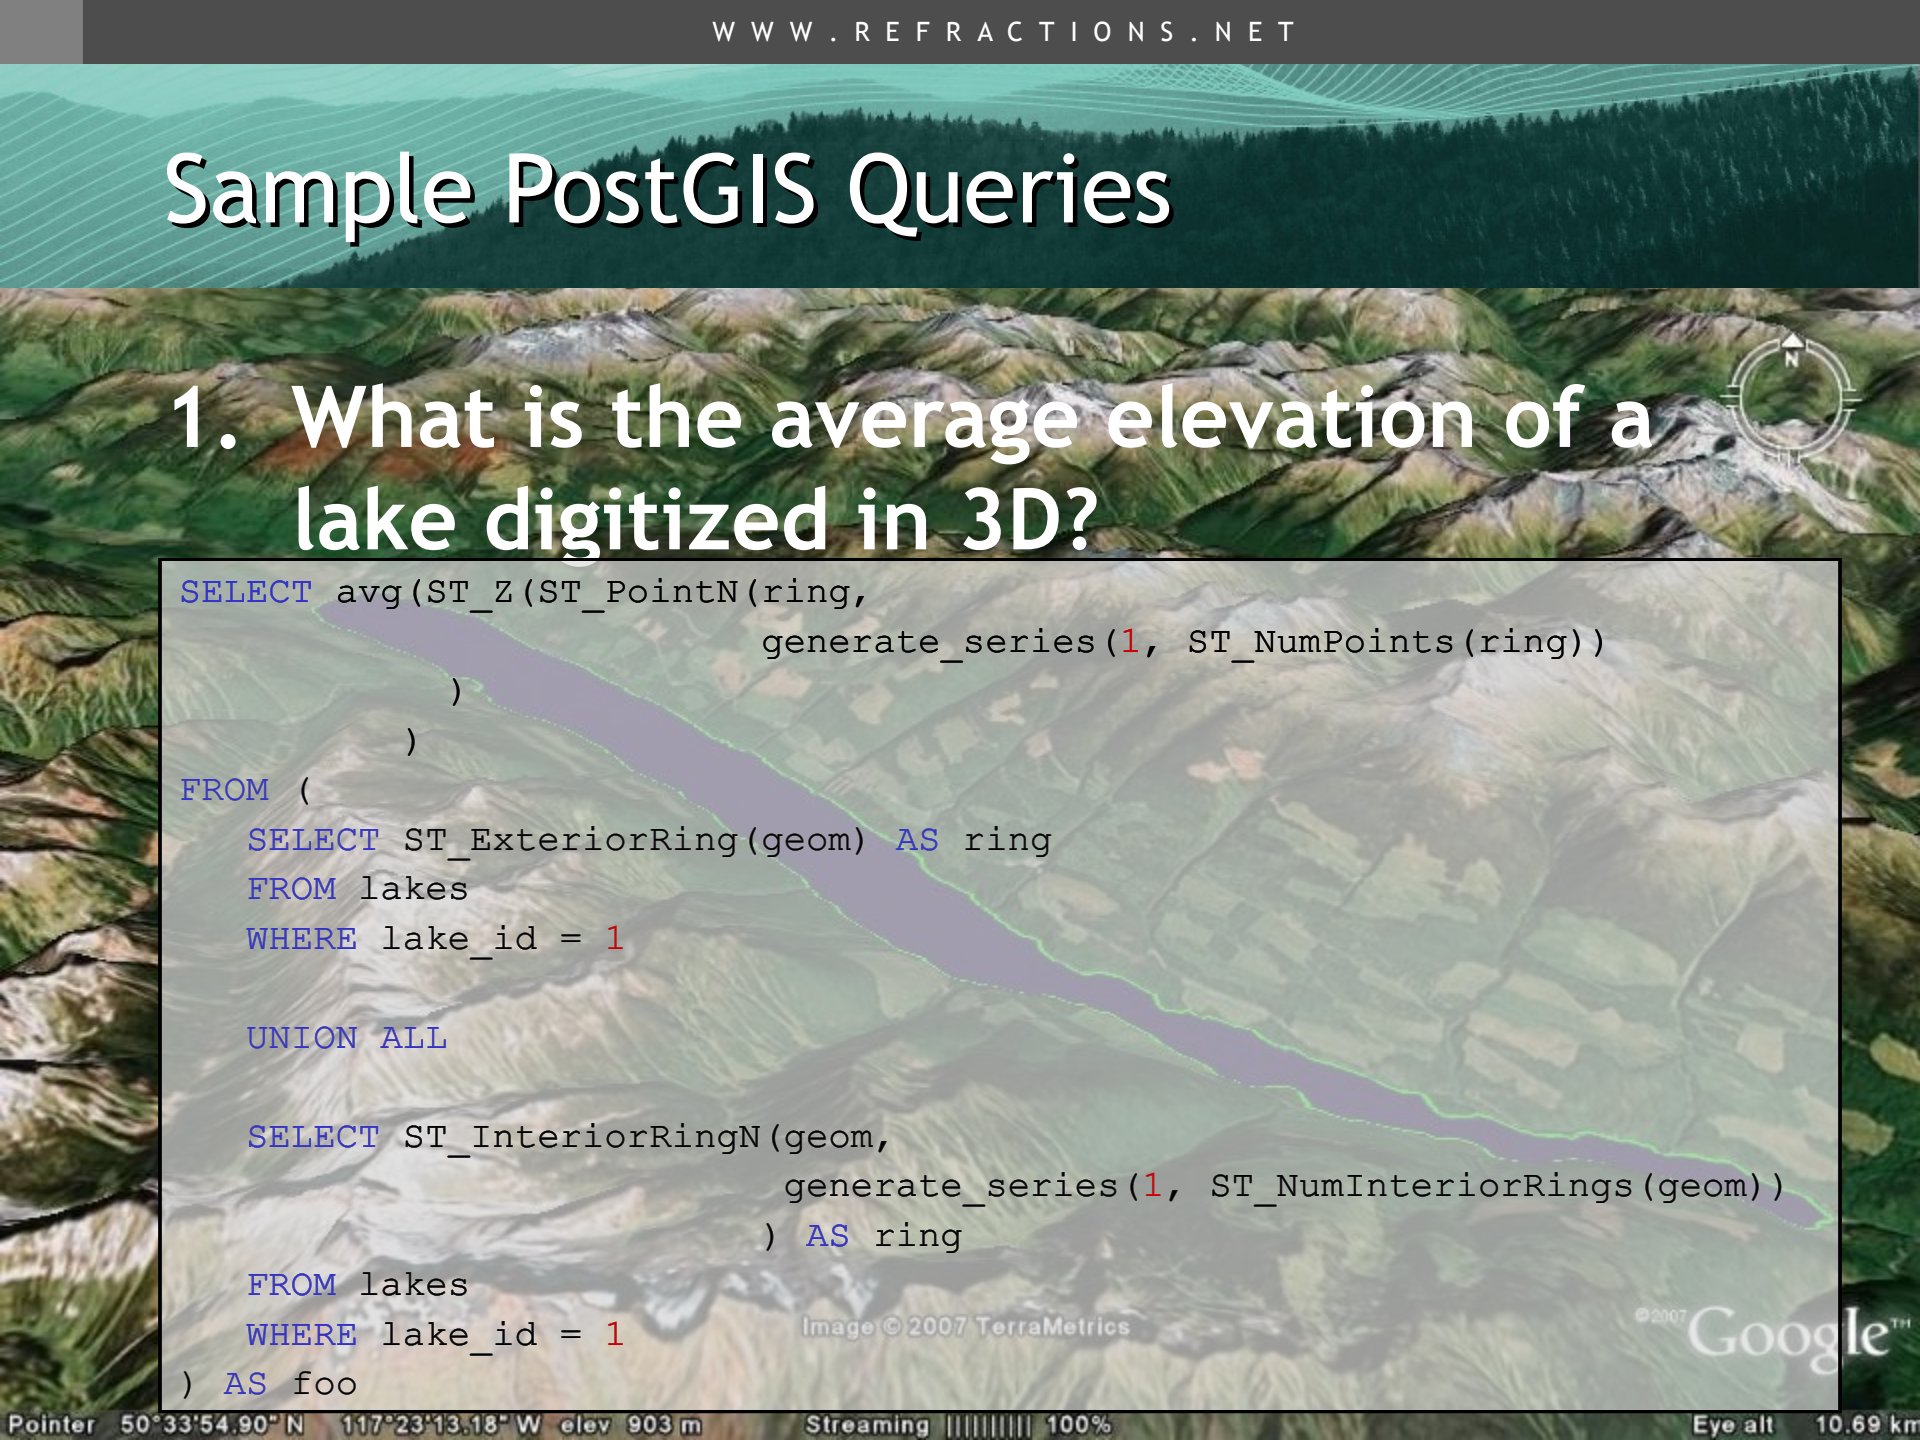
1. Identify the locations where clearcut logging occurs closer than 20m to a stream or river.



```
SELECT ST_Intersection(a.geom, ST_Buffer(b.geom, 20))  
FROM streams a, logging b  
WHERE ST_DWithin(a.geom, b.geom, 20)
```


Sample PostGIS Queries

1. What is the average elevation of a lake digitized in 3D?



```

SELECT avg(ST_Z(ST_PointN(ring,
                           generate_series(1, ST_NumPoints(ring))
                           )
          )
FROM (
  SELECT ST_ExteriorRing(geom) AS ring
  FROM lakes
  WHERE lake_id = 1

  UNION ALL

  SELECT ST_InteriorRingN(geom,
                           generate_series(1, ST_NumInteriorRings(geom))
                           ) AS ring
  FROM lakes
  WHERE lake_id = 1
) AS foo

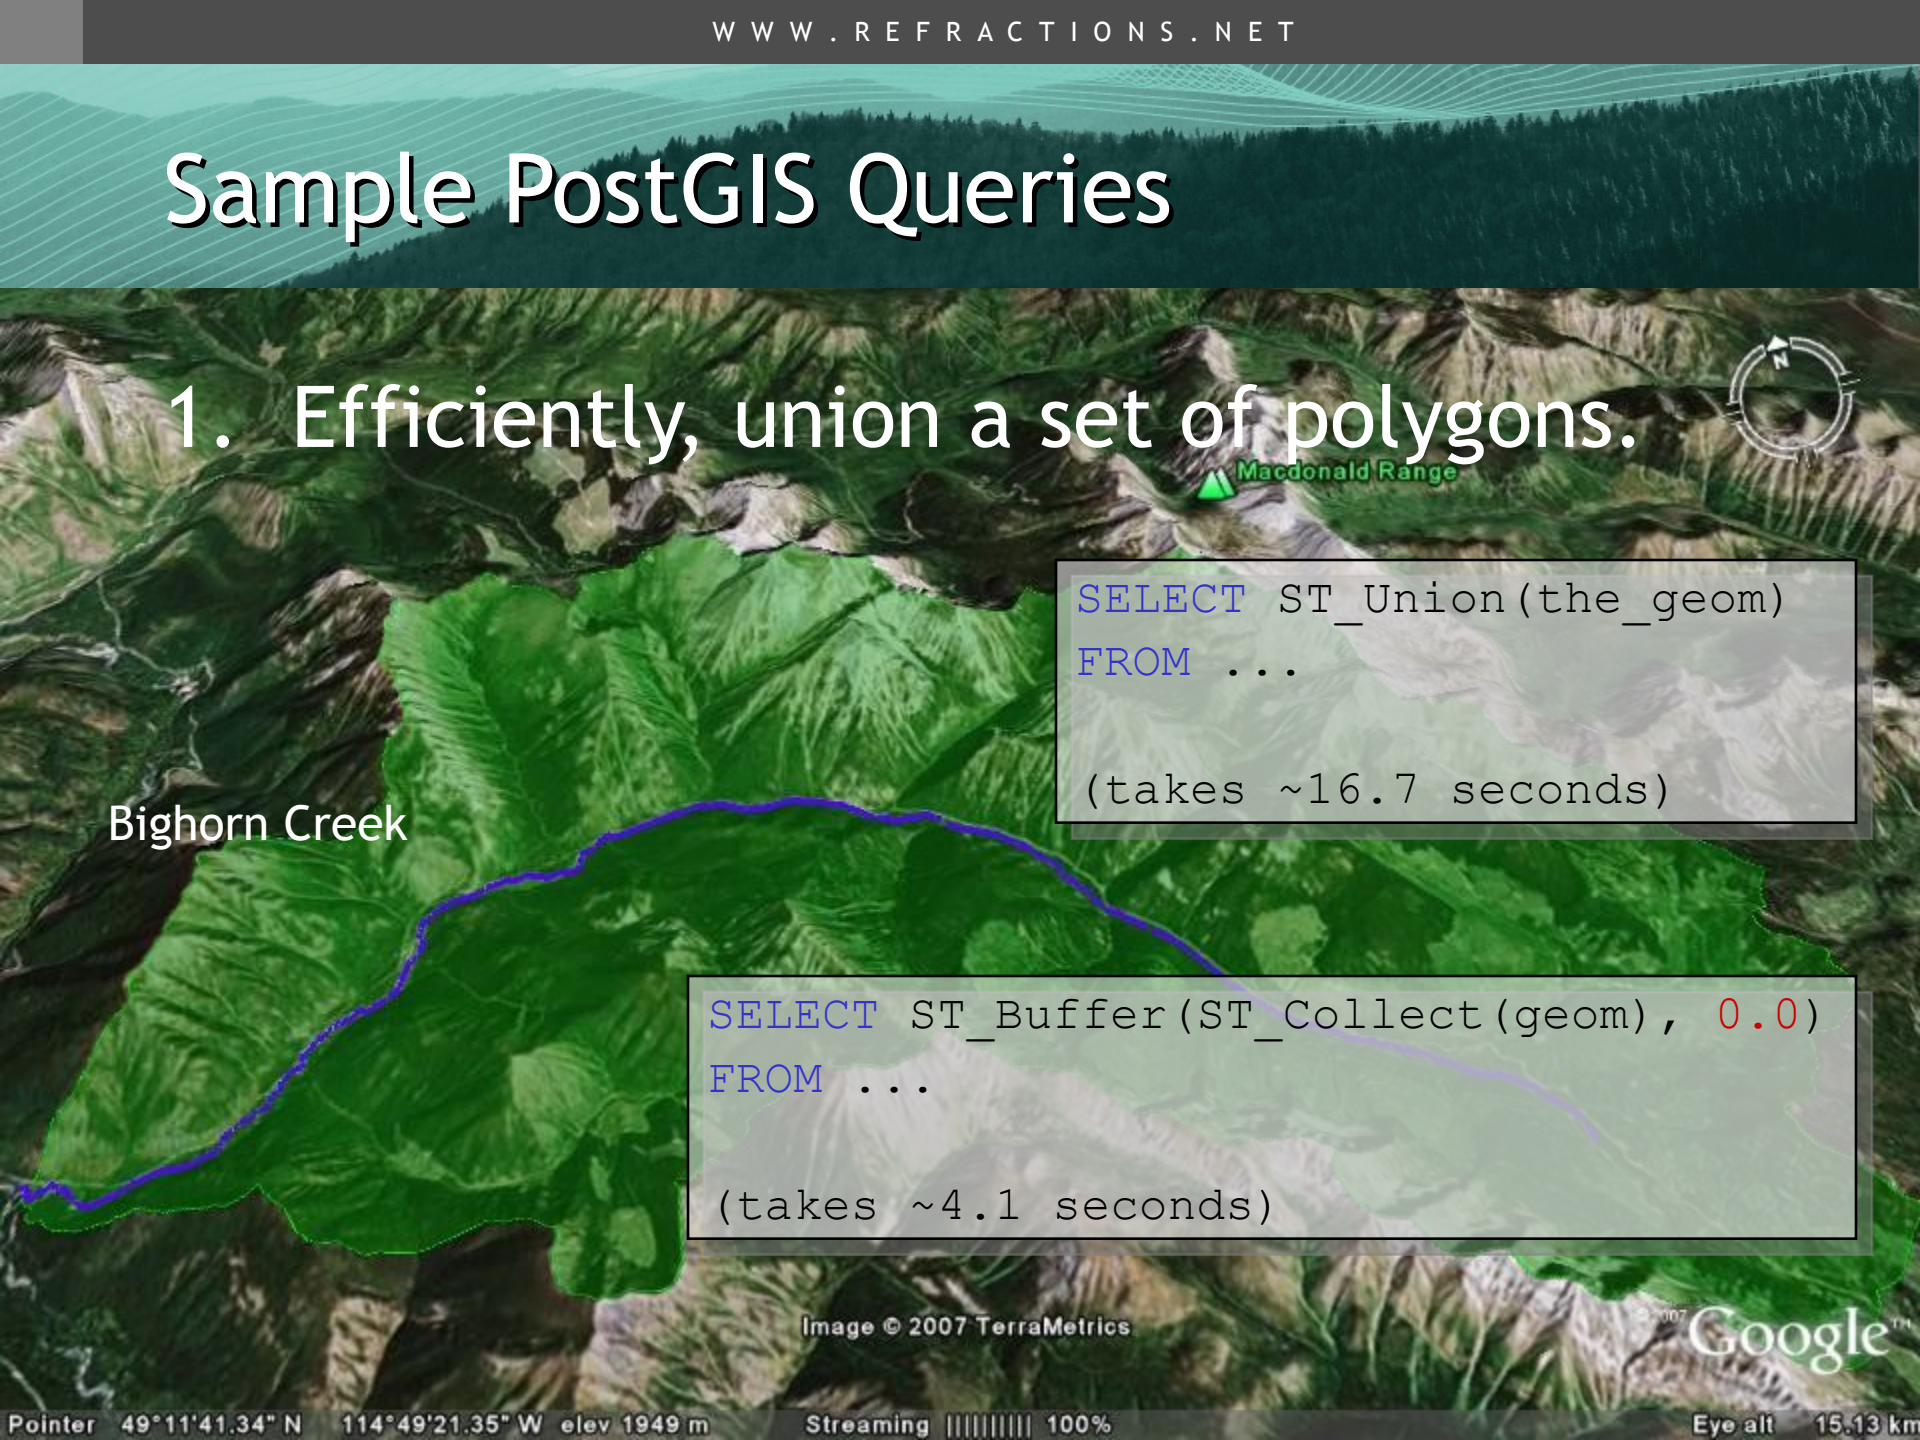
```

Image © 2007 TerraMetrics

© 2007 Google™

Sample PostGIS Queries

1. Efficiently, union a set of polygons.



```
SELECT ST_Union(the_geom)
FROM ...
```

(takes ~16.7 seconds)

Bighorn Creek

```
SELECT ST_Buffer(ST_Collect(geom), 0.0)
FROM ...
```


(takes ~4.1 seconds)

Image © 2007 TerraMetrics

Google™

Sample PostGIS Queries

1. Find all docks that are contained completely within a lake, not touching a lake bank.



What PostGIS functions would you use?

~~ST_Within?~~ ~~ST_Contains?~~ ~~ST_Touches?~~

DE-9IM The Dimensionally Extended - Nine Intersection Model

- Approach
 - make pair-wise tests of the intersections between the Interiors, Boundaries, and Exteriors of two geometries and to represent these relationships in an “intersection” matrix

DE-9IM The Dimensionally Extended - Nine Intersection Model

	Interior	Boundary	Exterior
Interior	$\dim(I(a) \cap I(b))$	$\dim(I(a) \cap B(b))$	$\dim(I(a) \cap E(b))$
Boundary	$\dim(B(a) \cap I(b))$	$\dim(B(a) \cap B(b))$	$\dim(B(a) \cap E(b))$
Exterior	$\dim(E(a) \cap I(b))$	$\dim(E(a) \cap B(b))$	$\dim(E(a) \cap E(b))$

Possible values:

{T, F, *, 0, 1, 2}

Where:

T == {0,1,2}

F == empty set

* == don't care

0 == dimensional 0 - point




1 == dimensional 1 - line

2 == dimensional 2 - area

DE-9IM The Dimensionally Extended - Nine Intersection Model

Geometry Topology




- Boundary
 - the set of geometries of the next lower dimension



DE-9IM The Dimensionally Extended - Nine Intersection Model

Geometry Topology

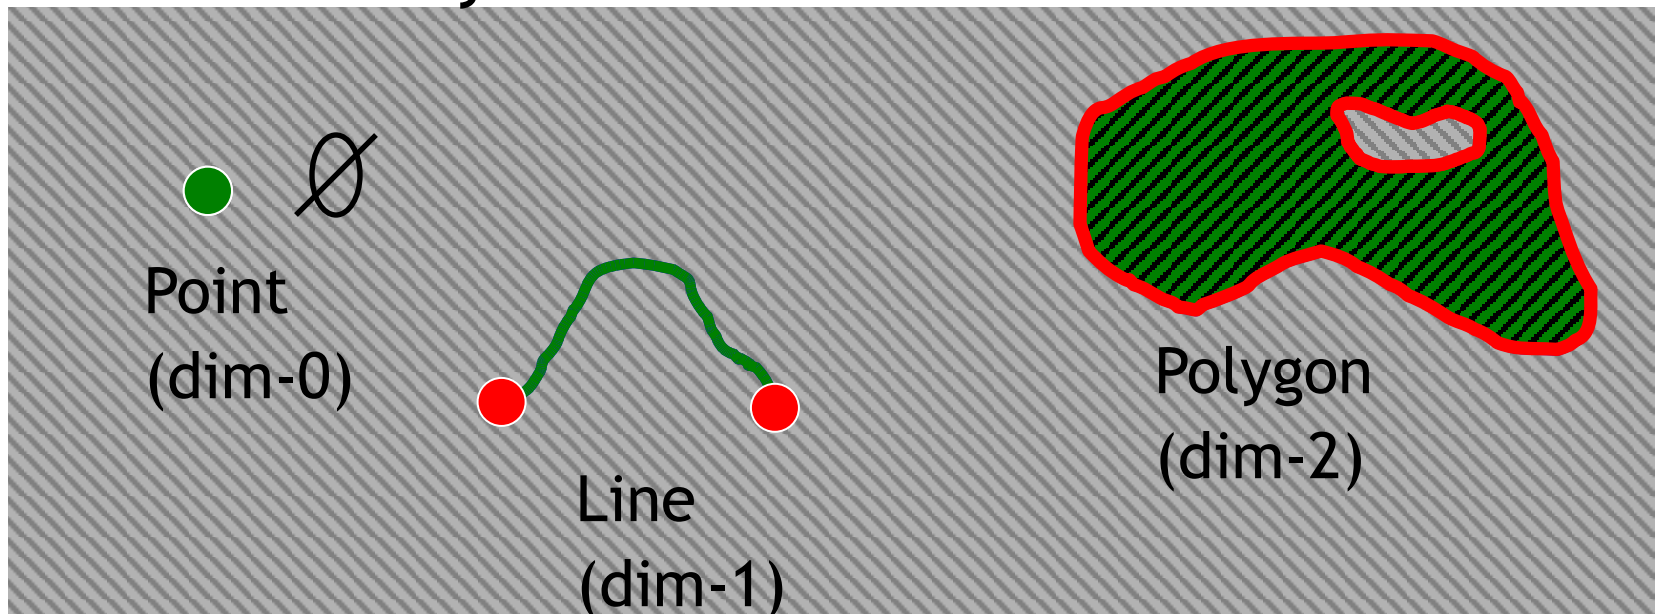
- Interior
 - the points that are left when the boundary points are removed




DE-9IM The Dimensionally Extended - Nine Intersection Model

Geometry Topology

- Exterior
 - consists of points not in the interior and boundary



DE-9IM The Dimensionally Extended - Nine Intersection Model




Interior Boundary Exterior

Interior	2	1	2
Boundary	1	0	1
Exterior	2	1	2

$ST_Relate(a, b) = '212101212'$

Sample PostGIS Queries


1. Find all docks that are contained completely within a lake, not touching a lake bank.



```
SELECT a.id  
FROM docks a, lakes b  
WHERE a.geom && b.geom  
AND ST_Relate(a.geom, b.geom, 'TFFTFF212');
```

Sample PostGIS Queries

1. Identify linear spatial features that intersect on a line and not at a point.



```
SELECT a.id, intersection(a.geom, b.geom)
FROM mylines a, mylines b
WHERE a.id != b.id
AND a.geom && b.geom
AND ST_Relate(a.geom, b.geom, '1*1***1**');
```

Table Inheritance

```
CREATE TABLE cities (  
  name      text,  
  population real,  
  altitude  int  
);
```

```
CREATE TABLE capitals (  
  province text  
) INHERITS (cities);
```

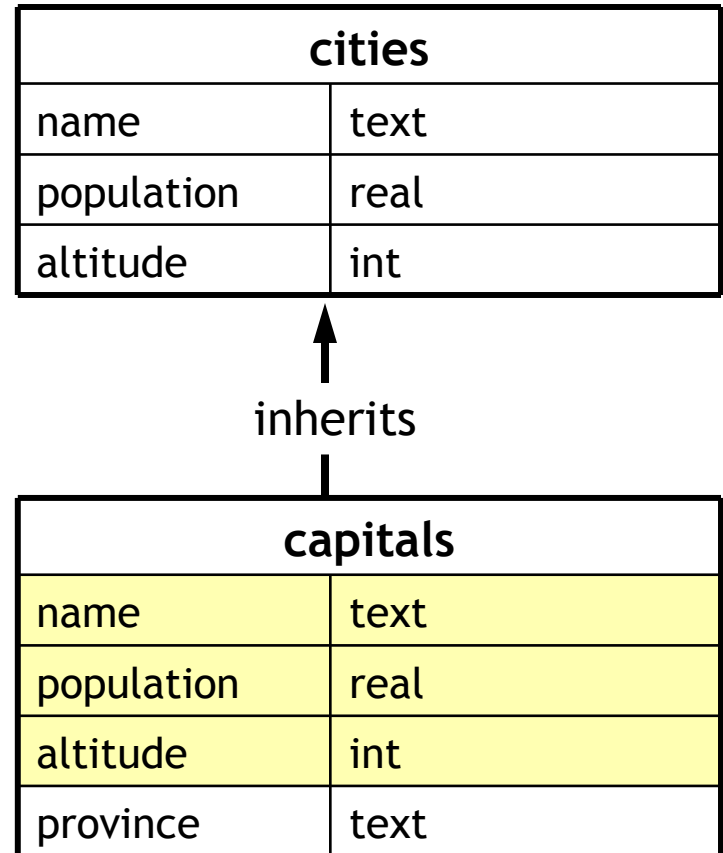


Table Inheritance

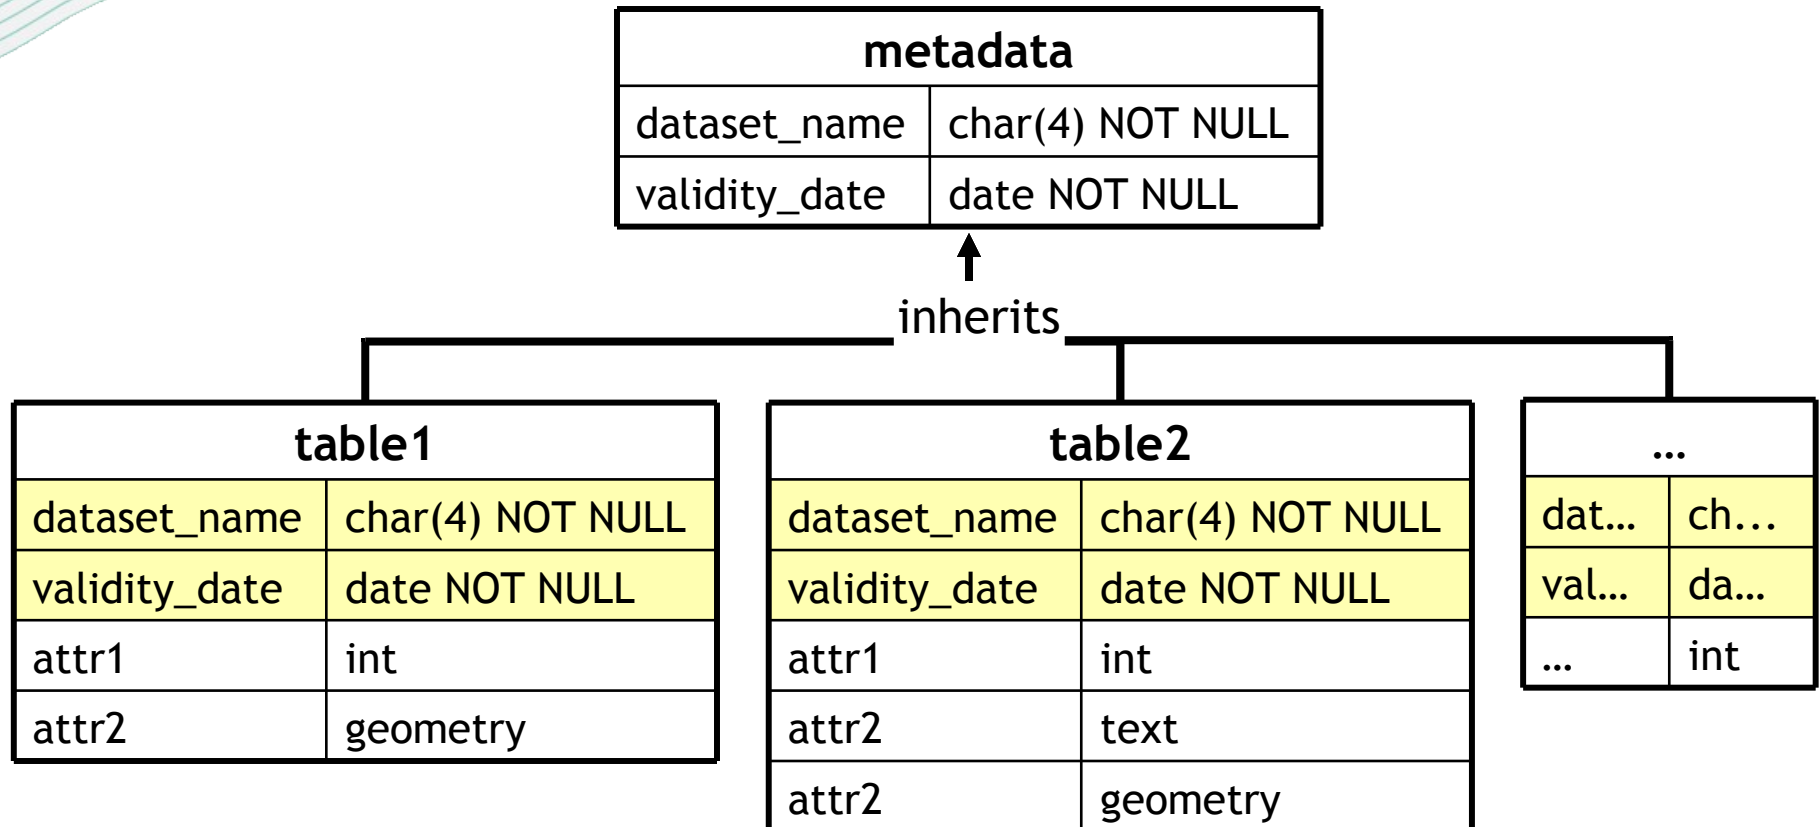


Table Partitioning

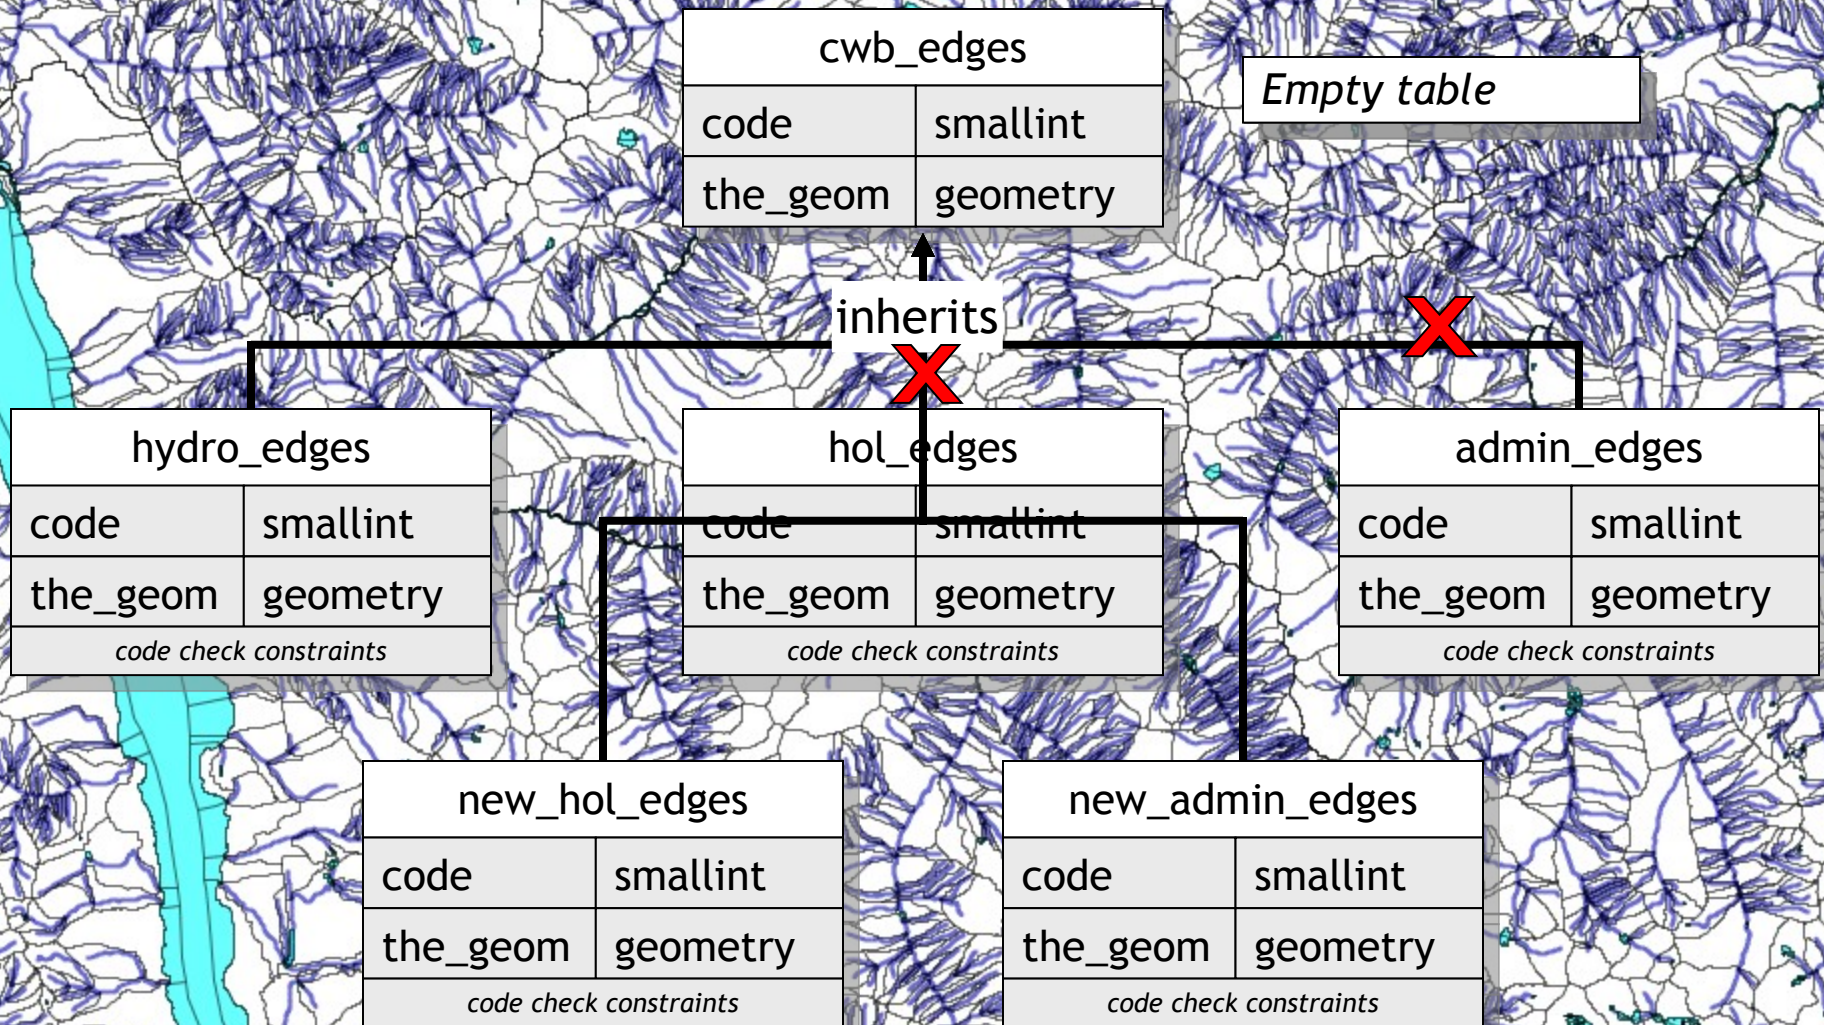


Table Partitioning

cwb_edges	
code	smallint
the_geom	geometry

hydro_edges	
code	smallint
the_geom	geometry
code check constraints	

```
-- ADD Rules to parent table
CREATE RULE insert_hydro AS
ON INSERT TO cwb_edges WHERE code = ...
DO INSTEAD
INSERT INTO hydro_edges (code, geom)
VALUES (NEW.code, NEW.geom);
```

```
CREATE TABLE hydro_edges (
) INHERITS cwb_edges;

ALTER TABLE hydro_edges
ADD CONSTRAINT code_check
CHECK (code = ...);
```

PostgreSQL Tuning

- The biggest bottleneck in a spatial database is I/O
- When setting up a server, invest in a:
 - *great* file system
 - RAID 5 - good for web servers, not spatial DBs
 - Recommend RAID 1/0
 - *good* memory
 - *adequate* CPU(s)

PostgreSQL Tuning

- postgresql.conf
 - Startup
 - checkpoint_segment_size
 - # of WAL files - 16MB each
 - Default: 3
 - Set to at least 10 or 30 for databases with heavy write activity or more for large database loads
 - Possibly store the xlog on a separate disk device
 - shared_buffers
 - Default: ~32MB
 - About 1/3 to 3/4 of available RAM

PostgreSQL Tuning

- postgresql.conf
 - Startup
 - constraint_exclusion
 - Default: "off"
 - Set to "on" to ensure the query planner will optimize as desired.

PostgreSQL Tuning

- postgresql.conf
 - Runtime
 - work_mem
 - Memory used for sort operations and complex queries
 - Default: 1MB
 - Adjust up for large dbs, complex queries, lots of RAM
 - Adjust down for many concurrent users or low RAM

PostgreSQL Tuning

- postgresql.conf
 - Runtime
 - maintenance_work_mem
 - Memory used for VACUUM, CREATE INDEX, etc.
 - Default:16MB
 - Generally too low - ties up I/O, locks objects while swapping memory.
 - Recommend 32MB to 256MB on production servers with lots of RAM, but depends on number of concurrent users.

PostgreSQL Tuning

- postgresql.conf
 - Runtime
 - On development systems with lots of RAM and few developers...

```
SET work_mem TO 1200000;
```

```
SET maintenance_work_mem TO 1200000;
```


PostgreSQL Tuning

- postgresql.conf
 - Runtime
 - client_min_messages

```
SET client_min_messages to DEBUG;
```

- Useful when writing PL/Pgsql functions.

```
CREATE FUNCTION my_function () RETURNS TEXT AS  
$BODY$  
BEGIN  
    ...  
    RAISE DEBUG 'myvar: %' var;  
    ...  
END
```

Performance Tips

- Spatial function calls can be expensive. Be efficient in their use - avoid unnecessary/duplicate function calls.
 - Use `St_Expand` where appropriate
 - Use one `relate` call instead of 2 or 3 other spatial calls.
 - Use `St_Distance()==0` instead of `intersects()` on large geometries
 - Avoid `St_Buffer()` unless you need a buffered geometry

Performance Tips

- Partition your data into Most Frequently Used (MFU) and Least Frequently Used (LFU).

Questions



Appendix A

```
// PostGIS and JTS
Class.forName("org.postgresql.Driver");
Connection conn =
    DriverManager.getConnection("jdbc:postgresql://...");

WKBReader wkbReader = new WKBReader();
WKBWriter wkbWriter = new WKBWriter();

String query =
    "SELECT the_geom FROM my_spatial_table
    WHERE the_geom && ST_GeomFromWKB(?, 3005)";
PreparedStatement pstmt = conn.prepareStatement(query);
pstmt.setBytes(1, wkbWriter.write(myJTSPolygon));

ResultSet rs = pstmt.executeQuery();
while(rs.next) {
    Geometry g = wkbReader.read(WKBReader.hexToBytes(
        rs.getString(1)));

    ...
    // Do stuff with Geometry
}
```

Appendix B

